

# Asynchronous Methods for Deep Reinforcement Learning

presented by Jason TOKO

# 背景与动机

- 传统经验认为，online RL方法与DNN相结合会导致不稳定。
- 主要原因：
  - 观察数据的不稳定性；
  - 样本间的相关性。
- 解决方法：
  - 使用Experience replay（经验回放）方法，可以减少不稳定性以及消除样本相关性。
- 限制：
  - 每次交互都需要大量的内存和计算；
    - 硬件要求较高，传统DRL方法依赖于GPU或者大型分布式架构等
  - 只能应用off-policy的方法进行学习。

# 异步RL框架

- 本文提出的异步RL框架，解决了经验回放存在的问题：
  - 1、异步地执行多个agent，通过并行的agent经历的不同状态，去除训练过程中产生样本之间的相关性；
  - 2、只需一个标准的多核CPU即可实现算法，在效果、时间和资源消耗上都优于传统方法；
  - 3、框架的通用性：不仅适用于off-policy、value-based方法，也适用于on-policy、policy-based方法，适用于离散和连续动作空间。
- 主要算法：
  - Asynchronous one-step Q-learning
  - Asynchronous one-step Sarsa
  - Asynchronous n-step Q-learning
  - Asynchronous advantage actor-critic

# Asynchronous one-step Q-learning

---

**Algorithm 1** Asynchronous one-step Q-learning - pseudocode for each actor-learner thread.

---

// Assume global shared  $\theta$ ,  $\theta^-$ , and counter  $T = 0$ . ← 全局共享参数 $\theta$ ,  $\theta^-$ ,  $T$

Initialize thread step counter  $t \leftarrow 0$

Initialize target network weights  $\theta^- \leftarrow \theta$

Initialize network gradients  $d\theta \leftarrow 0$

Get initial state  $s$

**repeat**

Take action  $a$  with  $\epsilon$ -greedy policy based on  $Q(s, a; \theta)$

Receive new state  $s'$  and reward  $r$

$$y = \begin{cases} r & \text{for terminal } s' \\ r + \gamma \max_{a'} Q(s', a'; \theta^-) & \text{for non-terminal } s' \end{cases}$$

Accumulate gradients wrt  $\theta$ :  $d\theta \leftarrow d\theta + \frac{\partial (y - Q(s, a; \theta))^2}{\partial \theta}$

$s = s'$

$T \leftarrow T + 1$  and  $t \leftarrow t + 1$

**if**  $T \bmod I_{target} == 0$  **then**

Update the target network  $\theta^- \leftarrow \theta$

**end if**

**if**  $t \bmod I_{AsyncUpdate} == 0$  or  $s$  is terminal **then**

Perform asynchronous update of  $\theta$  using  $d\theta$ .

Clear gradients  $d\theta \leftarrow 0$ .

**end if**

**until**  $T > T_{max}$

} 同步更新target network参数

} 异步更新main network参数

# Asynchronous one-step Sarsa

**Algorithm 1** Asynchronous one-step **Sarsa** - pseudocode for each actor-learner thread.

// Assume global shared  $\theta$ ,  $\theta^-$ , and counter  $T = 0$ . ← 全局共享参数 $\theta$ ,  $\theta^-$ ,  $T$

Initialize thread step counter  $t \leftarrow 0$

Initialize target network weights  $\theta^- \leftarrow \theta$

Initialize network gradients  $d\theta \leftarrow 0$

Get initial state  $s$

**repeat**

Take action  $a$  with  $\epsilon$ -greedy policy based on  $Q(s, a; \theta)$

Receive new state  $s'$  and reward  $r$

$$y = \begin{cases} r & \text{for terminal } s' \\ r + \gamma Q(s', a'; \theta^-) & \text{for non-terminal } s' \end{cases}$$

Accumulate gradients wrt  $\theta$ :  $d\theta \leftarrow d\theta + \frac{\partial(y - Q(s, a; \theta))^2}{\partial \theta}$

$s = s'$

$T \leftarrow T + 1$  and  $t \leftarrow t + 1$

**if**  $T \bmod I_{target} == 0$  **then**

Update the target network  $\theta^- \leftarrow \theta$

**end if**

**if**  $t \bmod I_{AsyncUpdate} == 0$  or  $s$  is terminal **then**

Perform asynchronous update of  $\theta$  using  $d\theta$ .

Clear gradients  $d\theta \leftarrow 0$ .

**end if**

**until**  $T > T_{max}$

} 同步更新target network参数

} 异步更新main network参数

# Asynchronous n-step Q-learning

- one-step方法中，每个r只作用于产生它的(s,a)的更新中，因而学习速度较慢
- 改进方法： n-step Q-learning:

$$TargetQ = r_t + \gamma r_{t+1} + \cdots + \gamma^{n-1} r_{t+n-1} + \gamma^n \max_{a'} Q(s', a'; \theta_i^-)$$

# Asynchronous n-step Q-learning

**Algorithm S2** Asynchronous n-step Q-learning - pseudocode for each actor-learner thread.

*// Assume global shared parameter vector  $\theta$ .*

*// Assume global shared target parameter vector  $\theta^-$ .*

*// Assume global shared counter  $T = 0$ .*

Initialize thread step counter  $t \leftarrow 1$

Initialize target network parameters  $\theta^- \leftarrow \theta$

Initialize thread-specific parameters  $\theta' = \theta$  ← 线程专用参数 $\theta'$

Initialize network gradients  $d\theta \leftarrow 0$

**repeat**

Clear gradients  $d\theta \leftarrow 0$

Synchronize thread-specific parameters  $\theta' = \theta$

$t_{start} = t$

Get state  $s_t$

**repeat**

Take action  $a_t$  according to the  $\epsilon$ -greedy policy based on  $Q(s_t, a; \theta')$

Receive reward  $r_t$  and new state  $s_{t+1}$

$t \leftarrow t + 1$

$T \leftarrow T + 1$

**until** terminal  $s_t$  **or**  $t - t_{start} == t_{max}$

$$R = \begin{cases} 0 & \text{for terminal } s_t \\ \max_a Q(s_t, a; \theta^-) & \text{for non-terminal } s_t \end{cases}$$

全局共享参数 $\theta$ ,  $\theta^-$ ,  $T$

产生n步序列

# Asynchronous n-step Q-learning

```
for  $i \in \{t - 1, \dots, t_{start}\}$  do  
     $R \leftarrow r_i + \gamma R$   
    Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \frac{\partial (R - Q(s_i, a_i; \theta'))^2}{\partial \theta'}$   
end for  
Perform asynchronous update of  $\theta$  using  $d\theta$ .  
if  $T \bmod I_{target} == 0$  then  
     $\theta^- \leftarrow \theta$   
end if  
until  $T > T_{max}$ 
```

累积n步梯度

异步更新main network参数

同步更新target network参数

---



# Asynchronous advantage actor-critic(A3C)

- Advantage function（优势函数）：

$$A(a_t, s_t) = Q(a_t, s_t) - V(s_t)$$

- 优势函数表现了动作  $a_t$  在状态  $s_t$  下的优劣程度~

- Advantage actor-critic:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

- 在A3C算法中，可用DNN去估计优势函数：

$$A(s_t, a_t; \theta', \theta'_v) = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta'_v) - V(s_t; \theta'_v)$$

- 实际中，引入熵项可避免过早收敛到次优解：

$$\nabla_{\theta'} \log \pi(a_t | s_t; \theta') (R_t - V(s_t; \theta_v)) + \beta \nabla_{\theta'} H(\pi(s_t; \theta'))$$

# Asynchronous advantage actor-critic(A3C)

**Algorithm S3** Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.

// Assume global shared parameter vectors  $\theta$  and  $\theta_v$  and global shared counter  $T = 0$

// Assume thread-specific parameter vectors  $\theta'$  and  $\theta'_v$

Initialize thread step counter  $t \leftarrow 1$

**repeat**

Reset gradients:  $d\theta \leftarrow 0$  and  $d\theta_v \leftarrow 0$ .

Synchronize thread-specific parameters  $\theta' = \theta$  and  $\theta'_v = \theta_v$

$t_{start} = t$

Get state  $s_t$

**repeat**

Perform  $a_t$  according to policy  $\pi(a_t|s_t; \theta')$

Receive reward  $r_t$  and new state  $s_{t+1}$

$t \leftarrow t + 1$

$T \leftarrow T + 1$

**until** terminal  $s_t$  **or**  $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$

**for**  $i \in \{t - 1, \dots, t_{start}\}$  **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

**end for**

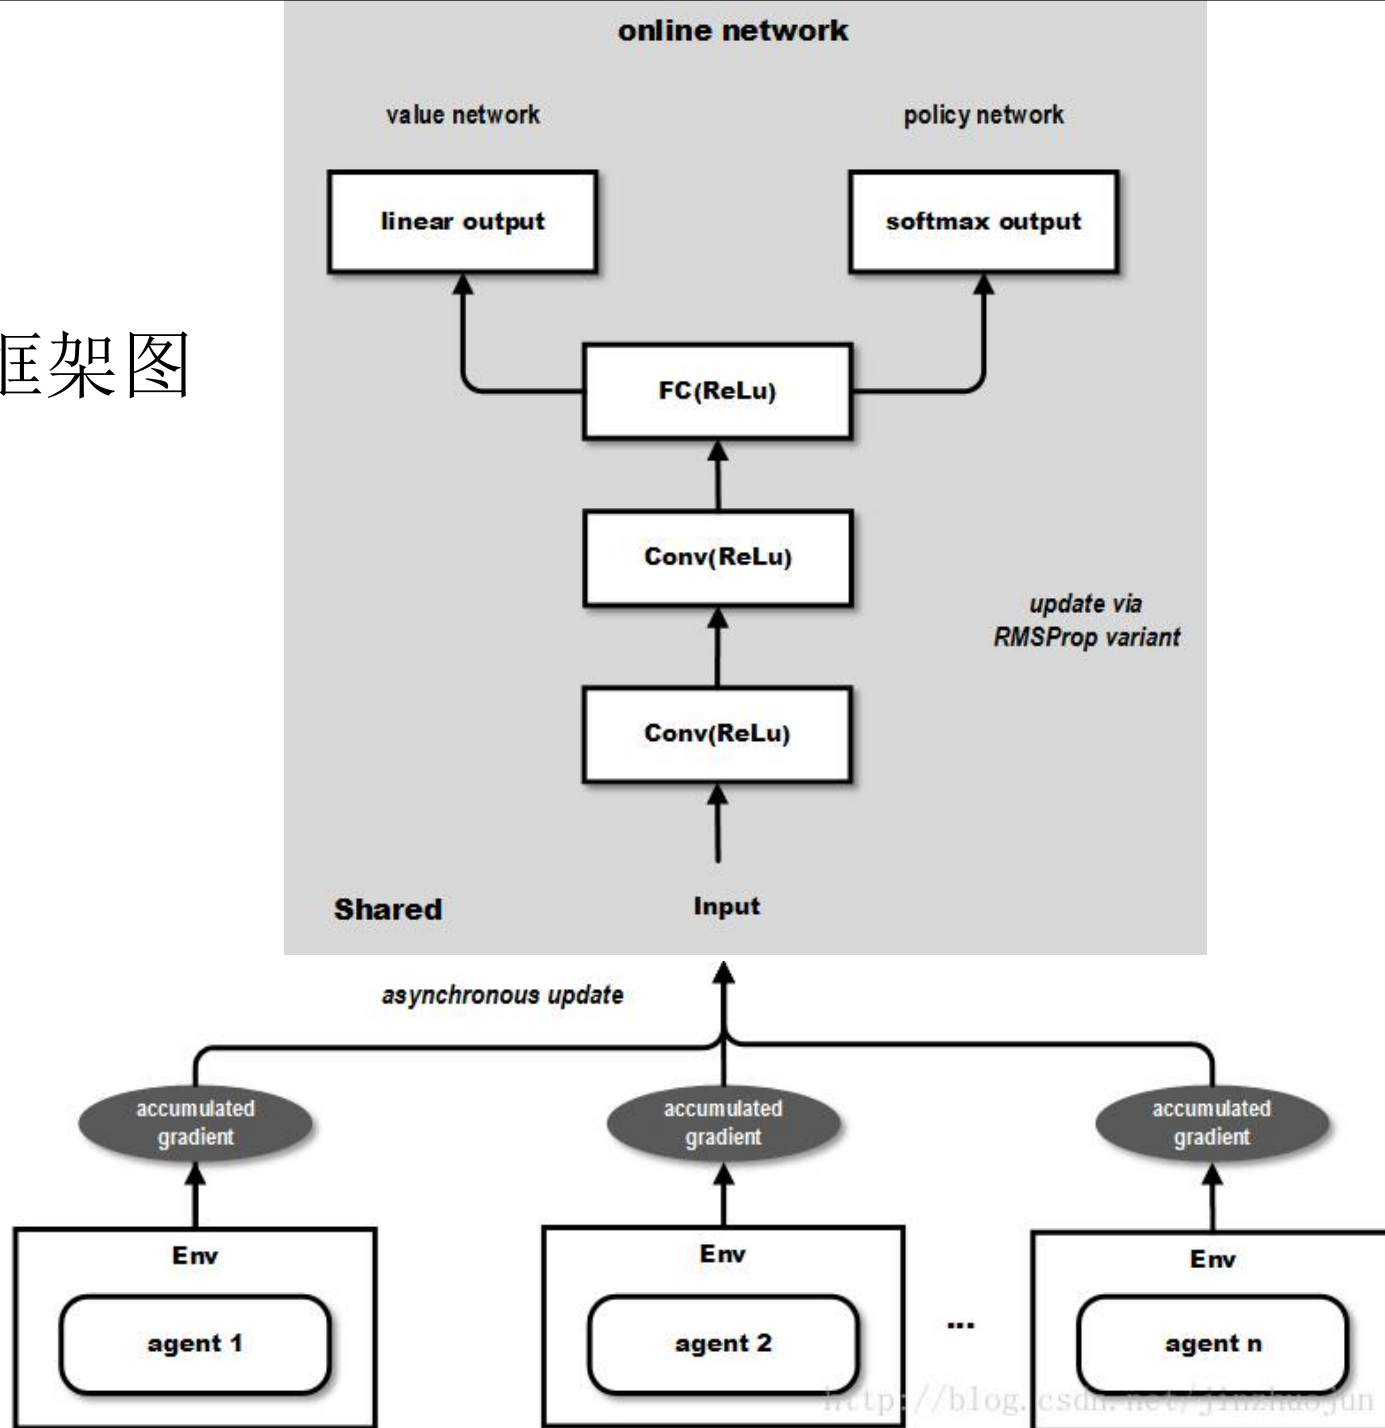
Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $\theta_v$  using  $d\theta_v$ .

**until**  $T > T_{max}$

产生n步序列

累积n步  
梯度

## A3C 框架图



# 总结

- 异步RL框架的着重点主要在于工程上的设计与优化。
- 异步框架可以解决DNN作为FA的不稳定性问题。
- 异步框架通用性很强，适用于on-policy和off-policy方法，适用于离散和连续动作空间。
- 我
- 随
- 便
- 填
- 充
- 一
- 下

- 引用:
- <http://blog.csdn.net/u013236946/article/details/73195035>
- <http://blog.csdn.net/jinzhuojun/article/details/72851548>